

```
import pandas as pd
import seaborn as sns
```

```
df = pd.read_csv("Dataset/Churn_Modelling.csv")
```

```
df.shape
```

```
(10000, 14)
```

```
df.columns
```

```
Index(['RowNumber', 'CustomerId', 'Surname', 'CreditScore',
       'Geography', 'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts',
       'HasCrCard', 'IsActiveMember', 'EstimatedSalary', 'Exited'],
      dtype='object')
```

```
df.head()
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age
0	1	15634602	Hargrave	619	France	Female	42
1	2	15647311	Hill	608	Spain	Female	41
2	3	15619304	Onio	502	France	Female	42
3	4	15701354	Boni	699	France	Female	39
4	5	15737888	Mitchell	850	Spain	Female	43

	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
0	2	0.00	1	1	1	
1	1	83807.86	1	0	1	
2	8	159660.80	3	1	0	
3	1	0.00	2	0	0	
4	2	125510.82	1	1	1	

	EstimatedSalary	Exited
0	101348.88	1
1	112542.58	0
2	113931.57	1
3	93826.63	0
4	79084.10	0

```
# Input Data
```

```
x =
```

```
df[['CreditScore', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard',
    'IsActiveMember', 'EstimatedSalary']]
```

```
# Output Data
```

```
y = df['Exited']
```


```
x
```

	CreditScore	Age	Tenure	Balance	NumOfProducts	HasCrCard	\
0	619	42	2	0.00	1	1	
1	608	41	1	83807.86	1	0	
2	502	42	8	159660.80	3	1	
3	699	39	1	0.00	2	0	
4	850	43	2	125510.82	1	1	
...	
9995	771	39	5	0.00	2	1	
9996	516	35	10	57369.61	1	1	
9997	709	36	7	0.00	1	0	
9998	772	42	3	75075.31	2	1	
9999	792	28	4	130142.79	1	1	

	IsActiveMember	EstimatedSalary
0	1	101348.88
1	1	112542.58
2	0	113931.57
3	0	93826.63
4	1	79084.10
...
9995	0	96270.64
9996	1	101699.77
9997	1	42085.58
9998	0	92888.52
9999	0	38190.78

```
[10000 rows x 8 columns]
```

```
sns.countplot(x = y);
```



```
y.value_counts()
0    7963
1    2037
Name: Exited, dtype: int64

from imblearn.over_sampling import RandomOverSampler
ros = RandomOverSampler(random_state=0)
x_res, y_res = ros.fit_resample(x, y)
y_res.value_counts()
1    7963
0    7963
Name: Exited, dtype: int64

# Normalize
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
x_scaled = scaler.fit_transform(x)
x_scaled
```

```
array([[ -0.32622142,  0.29351742, -1.04175968, ...,  0.64609167,
         0.97024255,  0.02188649],
       [ -0.44003595,  0.19816383, -1.38753759, ..., -1.54776799,
         0.97024255,  0.21653375],
       [ -1.53679418,  0.29351742,  1.03290776, ...,  0.64609167,
        -1.03067011,  0.2406869 ],
       ...,
       [  0.60498839, -0.27860412,  0.68712986, ..., -1.54776799,
         0.97024255, -1.00864308],
       [  1.25683526,  0.29351742, -0.69598177, ...,  0.64609167,
        -1.03067011, -0.12523071],
       [  1.46377078, -1.04143285, -0.35020386, ...,  0.64609167,
        -1.03067011, -1.07636976]])
```

Cross validation

```
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y,
random_state=0, test_size=0.25)
```

x.shape

```
(10000, 8)
```

x_test.shape

```
(2500, 8)
```

x_train.shape

```
(7500, 8)
```

```
from sklearn.neural_network import MLPClassifier
```

```
ann = MLPClassifier(hidden_layer_sizes=(100, 100, 100),
random_state=0, max_iter=100, activation='relu')
```

```
ann.fit(x_train, y_train)
```

```
C:\Users\ashre\anaconda3\envs\pandas\lib\site-packages\sklearn\
neural_network\_multilayer_perceptron.py:686: ConvergenceWarning:
Stochastic Optimizer: Maximum iterations (100) reached and the
optimization hasn't converged yet.
```

```
warnings.warn(
```

```
MLPClassifier(hidden_layer_sizes=(100, 100, 100), max_iter=100,
random_state=0)
```

```
y_pred = ann.predict(x_test)
```

```
from sklearn.metrics import ConfusionMatrixDisplay,
classification_report
from sklearn.metrics import accuracy_score
```

```
y_test.value_counts()
```


```
0    1991
```

```
1     509
```

```
Name: Exited, dtype: int64
```

```
ConfusionMatrixDisplay.from_predictions(y_test, y_pred)
```

```
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at  
0x1e13fcbde70>
```



```
accuracy_score(y_test,y_pred)
```

```
0.7556
```

```
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.80	0.93	0.86	1991
1	0.21	0.07	0.11	509
accuracy			0.76	2500
macro avg	0.50	0.50	0.48	2500

weighted avg	0.68	0.76	0.71	2500
--------------	------	------	------	------