

```
import pandas as pd
```

```
df = pd.read_csv('Dataset/emails.csv')
```

```
df.shape
```

```
(5172, 3002)
```

```
df.head()
```

	Email No.	the	to	ect	and	for	of	a	you	hou	...	connevey
0	Email 1	0	0	1	0	0	0	2	0	0	...	0
1	Email 2	8	13	24	6	6	2	102	1	27	...	0
2	Email 3	0	0	1	0	0	0	8	0	0	...	0
3	Email 4	0	5	22	0	5	1	51	2	10	...	0
4	Email 5	7	6	17	1	5	2	57	0	9	...	0

	valued	lay	infrastructure	military	allowing	ff	dry
0	0	0		0	0	0	0
1	0	0		0	0	0	1
2	0	0		0	0	0	0
3	0	0		0	0	0	0
4	0	0		0	0	0	1

```
[5 rows x 3002 columns]
```

```
# Input Data
```

```
x = df.drop(['Email No.', 'Prediction'], axis = 1)
```

```
# Output Data
```

```
y = df['Prediction']
```

```
x.shape
```

```
(5172, 3000)
```

```
x.dtypes
```


the	int64
to	int64

```
ect          int64
and          int64
for          int64
...
infrastructure int64
military      int64
allowing      int64
ff           int64
dry          int64
Length: 3000, dtype: object
```

```
set(x.dtypes)
```

```
{dtype('int64')}
```

```
import seaborn as sns
sns.countplot(x = y);
```



```
y.value_counts()
```

```
0    3672
```

```
1    1500
```

```
Name: Prediction, dtype: int64
```

```

# Feature Scaling
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
x_scaled = scaler.fit_transform(x)

x_scaled
array([[0.          , 0.          , 0.          , ..., 0.          , 0.
        ,
        0.          ],
       [0.03809524, 0.09848485, 0.06705539, ..., 0.          ,
0.00877193,
        0.          ],
       [0.          , 0.          , 0.          , ..., 0.          , 0.
        ,
        0.          ],
       ...,
       [0.          , 0.          , 0.          , ..., 0.          , 0.
        ,
        0.          ],
       [0.00952381, 0.0530303 , 0.          , ..., 0.          ,
0.00877193,
        0.          ],
       [0.1047619 , 0.18181818, 0.01166181, ..., 0.          , 0.
        ,
        0.          ]])

# Cross Validation
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x_scaled, y,
random_state = 0 , test_size = 0.25)

x_scaled.shape
(5172, 3000)

x_train.shape
(3879, 3000)

x_test.shape
(1293, 3000)

# Import the class
from sklearn.neighbors import KNeighborsClassifier

# Create the object
knn = KNeighborsClassifier(n_neighbors=5)

# Train the algorithm
knn.fit(x_train, y_train)

```

```

KNeighborsClassifier()


# predict on test data
y_pred = knn.predict(x_test)

# import the evaluation metrics
from sklearn.metrics import ConfusionMatrixDisplay, accuracy_score
from sklearn.metrics import classification_report

ConfusionMatrixDisplay.from_predictions(y_test, y_pred)

<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at
0x21693f3a0e0>

```



```

y_test.value_counts()

0    929
1    364
Name: Prediction, dtype: int64

accuracy_score(y_test, y_pred)

0.871616395978345

print(classification_report(y_test, y_pred))

```

	precision	recall	f1-score	support
0	0.98	0.84	0.90	929
1	0.70	0.95	0.81	364
accuracy			0.87	1293
macro avg	0.84	0.89	0.85	1293
weighted avg	0.90	0.87	0.88	1293

```
import numpy as np
import matplotlib as plt
```

```
error = []
for k in range(1,41):
    knn = KNeighborsClassifier(n_neighbors=k)
    knn.fit(x_train, y_train)
    pred = knn.predict(x_test)
    error.append(np.mean(pred != y_test))
```

error

```
[0.10827532869296211,
 0.10982211910286156,
 0.12296983758700696,
 0.11523588553750967,
 0.12838360402165508,
 0.1214230471771075,
 0.15158546017014696,
 0.14849187935034802,
 0.17246713070378963,
 0.16705336426914152,
 0.1871616395978345,
 0.18329466357308585,
 0.21500386697602475,
 0.21345707656612528,
 0.22815158546017014,
 0.2266047950502707,
 0.23588553750966745,
 0.23356535189481825,
 0.2459396751740139,
 0.24361948955916474,
 0.2559938128383604,
 0.2552204176334107,
 0.2699149265274555,
 0.2691415313225058,
 0.2822892498066512,
 0.28306264501160094,
 0.2954369682907966,
 0.2923433874709977,
```

```
0.3039443155452436,  
0.300077339520495,  
0.30549110595514306,  
0.30549110595514306,  
0.31245166279969067,  
0.31245166279969067,  
0.3194122196442382,  
0.317092034029389,  
0.32637277648878577,  
0.32559938128383603,  
0.33410672853828305,  
0.3325599381283836]
```

```
knn = KNeighborsClassifier(n_neighbors=1)  
knn.fit(x_train, y_train)
```

```
KNeighborsClassifier(n_neighbors=1)
```

```
y_pred = knn.predict(x_test)  
accuracy_score(y_test, y_pred)
```

```
0.8917246713070379
```

```
from sklearn.svm import SVC  
svm = SVC(kernel='poly')  
svm.fit(x_train, y_train)
```

```
SVC(kernel='poly')
```

```
y_pred = svm.predict(x_test)  
accuracy_score(y_test, y_pred)
```

```
0.7548337200309359
```

```
# Linear:0.9767981438515081  
# RBF:0.9450889404485692  
# Poly:0.7548337200309359
```